

Hill Climbing

In numerical analysis, hill climbing is a mathematical optimization technique which belongs to the family of local search. It is an iterative algorithm that starts with an arbitrary solution to a problem, then attempts to find a better solution by making an incremental change to the solution. If the change produces a better solution, another incremental change is made to the new solution, and so on until no further improvements can be found.

For example, hill climbing can be applied to the travelling salesman problem. It is easy to find an initial solution that visits all the cities but will likely be very poor compared to the optimal solution. The algorithm starts with such a solution and makes small improvements to it, such as switching the order in which two cities are visited. Eventually, a much shorter route is likely to be obtained.

Hill climbing finds optimal solutions for convex problems - for other problems it will find only local optima (solutions that cannot be improved upon by any neighbouring configurations), which are not necessarily the best possible solution (the global optimum) out of all possible solutions (the search space). Examples of algorithms that solve convex problems by hill-climbing include the simplex algorithm for linear programming and binary search. To attempt to avoid getting stuck in local optima, one could use restarts (i.e. repeated local search), or more complex schemes based on iterations (like iterated local search), or on memory (like reactive search optimization and tabs search), or on memory-less stochastic modifications (like simulated annealing).

Pseudocode:

Algorithm Discrete Space Hill Climbing	Algorithm Continuous Space Hill Climbing is
<pre>currentNode := startNode loop do L := NEIGHBORS(currentNode) nextEval := -INF nextNode := NULL for all x in L do if EVAL(x) > nextEval then nextNode := x nextEval := EVAL(x) if nextEval ≤ EVAL(currentNode) then // Return current node since no better neighbors exist return currentNode currentNode := nextNode</pre>	<pre>currentPoint := initialPoint // the zero-magnitude vector is common stepSize := initialStepSizes // a vector of all 1's is common acceleration := someAcceleration // a value such as 1.2 is common candidate[0] := -acceleration candidate[1] := -1 / acceleration candidate[2] := 1 / acceleration candidate[3] := acceleration bestScore := EVAL(currentPoint) loop do beforeScore := bestScore for each element i in currentPoint do beforePoint := currentPoint[i] bestStep := 0 for j from 0 to 3 do // try each of 4 candidate locations step := stepSize[i] × candidate[j] currentPoint[i] := beforePoint + step score := EVAL(currentPoint) if score > bestScore then bestScore := score bestStep := step if bestStep is 0 then currentPoint[i] := beforePoint stepSize[i] := stepSize[i] / acceleration else currentPoint[i] := beforePoint + bestStep stepSize[i] := bestStep // accelerate if (bestScore - beforeScore) < epsilon then return currentPoint</pre>